


The Impact of Inclement Weather and Subway Outages on Citi Bike Ridership

Brian Weinstein
Andy Enkeboll
Ashutosh Nanda

<https://github.com/BrianWeinstein/citibike-predictions>

Overview



Data Source: Subway Status v1

Alert Archive

[Get Free Service Alerts](#)

Start:  End: 

	Sent Date	Agency	Subject	Message
>	5/14/2015 7:12:16 PM	Subway	MANH, 4 and 5 Trains, Signal Problems	b/d, 4 and 5 trains delayed, some s/b 5 via 2 from 149 St to Nevins St due to signal problems at 86 St. Allow additional travel time.
>	5/14/2015 7:00:24 PM	LIRR	Port Washington Br./Great Neck W/B Train Operating 13 Minutes Late	The 6:44PM train from Great Neck due Penn at 7:15PM is operating 13 minutes late due to switch trouble.
>	5/14/2015 6:58:54 PM	LIRR	Port Washington Westbound Train 11 Minutes Late	The 6:24PM train from Port Washington due Penn at 7:01PM is operating 11 minutes late due to congestion on the branch caused by switch trouble at Great Neck.
>	5/14/2015 6:48:47 PM	LIRR	Babylon Eastbound Train 11 Minutes Late	The 5:59PM train from Penn due Babylon at 7:01PM is operating 11 minutes late due to a late train ahead.


Data Source: Subway Status v2




Status and delays on the Wednesday May 13, 2015

The overall status of the 1 Train is **Planned Work** for 22 % of the time, **Good Service** for 32 % of the time, **Service Change** for 27 % of the time, **Delays** for 19 % of the time.

Today's Status Evolution



Today's Status Breakdown



```
var yourLabels = [ "",  
    "Good Service",  
    "Delays",  
    "Planned Work",  
    "Service Change",  
    "Suspended" ];  
  
series: [{  
    name: 'Status',  
    data: [3,3,3,3,3,3,3,3,3,3,3,3,  
        3,3,3,3,3,3,3,3,3,3,1,1,  
        1,1,1,1,1,4,2,1,1,1,1,1,  
        1,1,4,4,4,4,4,4,4,4,4,4,  
        4,4,4,4,4,4,4,4,4,4,4,4,  
        4,4,4,2,1,1,1,1,1,1,1,2,  
        2,2,2,2,2,2,2,2,2,2,2,2,  
        2,2,2,1,1,1,1,1,1,1,1,1,  
        1]  
}]
```

Data Source: CitiBike Trips


Citi Bike Trip Histories


Below are links to downloadable files of Citi Bike trip data. The data includes:

- Trip Duration (seconds)
- Start Time and Date
- Stop Time and Date
- Start Station Name
- End Station Name
- Station ID
- Station Lat/Long
- Bike ID
- User Type (Customer = 24-hour pass or 7-day pass user; Subscriber = Annual Member)
- Gender (Zero=unknown; 1=male; 2=female)
- Year of Birth

[July 2013](#) (25.8 MB)
[August 2013](#) (30.6 MB)
[September 2013](#) (31.6 MB)
[October 2013](#) (31.5 MB)
[November 2013](#) (20.6 MB)
[December 2013](#) (13.6 MB)
[January 2014](#) (9.2 MB)
[February 2014](#) (6.9 MB)
[March 2014](#) (13.4 MB)
[April 2014](#) (20.4 MB)
[May 2014](#) (26.3 MB)
[June 2014](#) (28.5 MB)
[July 2014](#) (29.4 MB)
[August 2014](#) (29.2 MB)
[September 2014](#) (28.8 MB)
[October 2014](#) (24.9 MB)
[November 2014](#) (16 MB)
[December 2014](#) (12.1 MB)

Data Source: New York Weather

**NOAA** NATIONAL CLIMATIC DATA CENTER
NATIONAL OCEANIC AND ATMOSPHERIC ADMINISTRATION

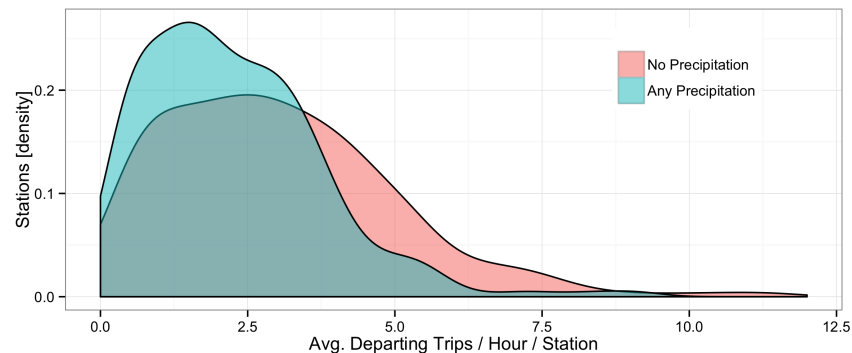

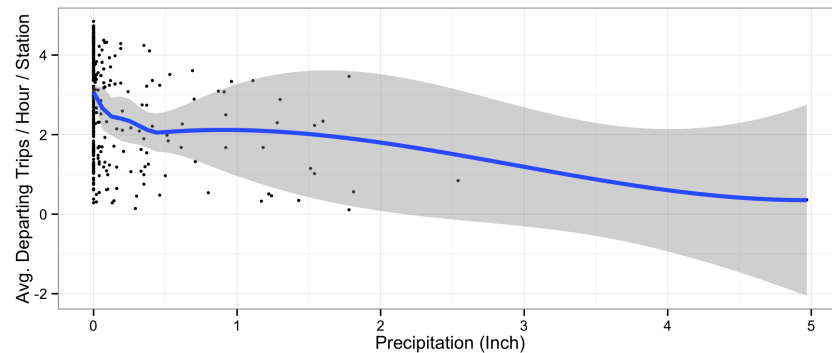

Your data request is complete and ready to download
<http://www1.ncdc.noaa.gov/pub/orders/cdo/526965.csv>

Order Summary
Order Number: [526965](#)
Order Status: Complete
Order Format: Custom GHCN-Daily CSV

Requested Data
Custom Options: Station name, Geographic location
Stations: **GHCND:USW00094728** - NEW YORK CENTRAL PARK OBS
BELVEDERE TOWER, NY US
Data Types: **TMAX** - Maximum temperature (tenths of degrees C)
SNWD - Snow depth (mm)
SNOW - Snowfall (mm)
PRCP - Precipitation (tenths of mm)
TMIN - Minimum temperature (tenths of degrees C)



Data Exploration





Prediction Model

- Tested 4 regression models on two subsets of predictors:

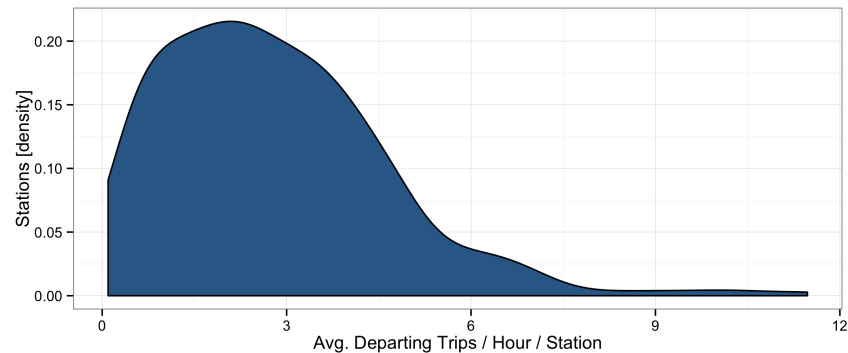
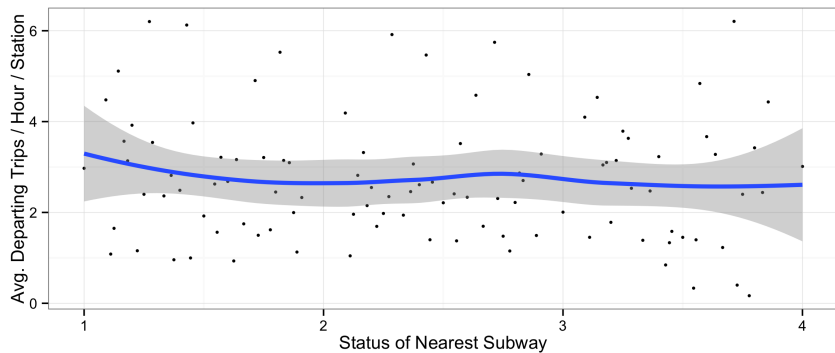

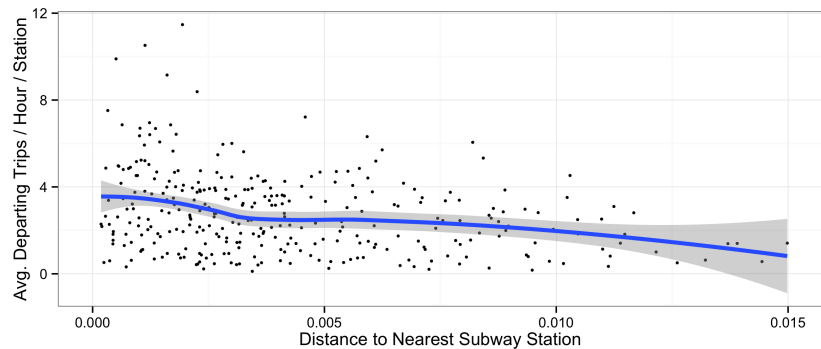
Subset 1: `trips ~ hour + weekday + nearestSubStationDist + avgSubStationStatus + anyPrecip + maxTemp`

Subset 2: `trips ~ hour + weekday + citiStationID + avgSubStationStatus + anyPrecip + maxTemp`

<i>Test RMSE</i>	Ridge	Lasso	Linear Least Squares	Polynomial Least Squares
Subset 1	3.2910	3.2885	3.2884	3.2811
Subset 2	2.9843	2.9808	2.9809	2.9727



Additional: Data Exploration



Additional: Optimization

When using `set.seed(12)`, Lasso didn't shrink any of the coefficients to 0, although it does with other seeds.

```
>coef(lasso.model, s=lasso.bestlambda)
(Intercept)      -3.29118      weekday1      0.67272
hour1           -0.34314      citiStationID127    3.56446
hour2           -0.43761      citiStationID143   -0.19339
hour3           -0.48536      citiStationID167    2.74610
hour4           -0.47186      citiStationID228    1.33829
hour5           -0.40206      citiStationID254    0.27684
hour6            0.22949      citiStationID259    0.85215
hour7            1.31776      citiStationID294    3.27361
hour8            3.65633      citiStationID298   -0.70295
hour9            3.08644      citiStationID331    0.13698
hour10           1.89055      citiStationID332    0.27645
hour11           2.08708      citiStationID339   -0.58213
hour12           2.42439      citiStationID352    2.22051
hour13           2.56377      citiStationID357    1.69681
hour14           2.78327      citiStationID389   -0.41492
hour15           2.82930      citiStationID430    0.81722
hour16            3.38743      citiStationID431   -0.81171
hour17            4.56809      citiStationID485    1.70873
hour18            4.35563      citiStationID494    3.09902
hour19           2.80921      citiStationID536    2.79051
hour20           1.68588      avgSubStationStatus -0.11149
hour21           0.91142      anyPrecip1        -0.63762
hour22           0.57315      maxTemp            0.04402
hour23           0.21723
```

Data Source: MTA Subway Locations

```
curl https://data.ny.gov/api/views/i9wp-a4ja/rows.csv | \
  cut -d , -f 2-16 > datasets/mta_station_data_raw.csv

(head -n 1 datasets/mta_station_data_raw.csv && \
  tail -n +2 datasets/mta_station_data_raw.csv | \
  sort | uniq) > datasets/mta_station_data.csv

rm datasets/mta_station_data_raw.csv

awk -F',' 'NR==1 {print "Station ID","$2","$3","$4"} NR>1 {print NR-1,""$1": "$2","$3","$4} ' \
  OFS=, datasets/mta_station_data.csv > datasets/mta_station_location.csv

awk -F',' 'NR==1 {print "Station ID,Line"} NR>1 {for (i=5; i<=15; i++) {if (length($i) > 0) {print NR-1,""$i}}}' \
  OFS=, datasets/mta_station_data.csv > datasets/mta_station_lines.csv

rm datasets/mta_station_data.csv
```

Data Source: Subway Status v1

```
# initialize webdriver instance and visit url
url = 'http://archive.mymtaalerts.com/messagearchive.aspx'
browser = webdriver.Firefox()
browser.get(url)

# set start date
start_date = browser.find_element_by_id('RadDatePickerStart_dateInput')
start_date.send_keys(Keys.COMMAND, 'a')
start_date.send_keys('1/1/2014')
start_date.send_keys(Keys.ENTER)

# set end date
end_date = browser.find_element_by_id('RadDatePickerEnd_dateInput')
end_date.send_keys(Keys.COMMAND, 'a')
end_date.send_keys('12/31/2014')
end_date.send_keys(Keys.ENTER)
```

Data Source: Subway Status v2

```
while pull_date < max_date:
    for line in lines:
        url = base_url.format(line=line, date=pull_date.isoformat())
        try:
            page = urllib2.urlopen(url)
            soup = BeautifulSoup(page.read())
            # this is super brittle, but this site is very well organized so it works
            script = soup.find_all('script')[10]
            data = script.text.split("data: ")[1].split("\r\n")[0].replace("null", "None")
            data = ast.literal_eval(data)

            prev_d = None
            for i in range(len(data)):
                if data[i] is None:
                    data[i] = prev_d if prev_d else 1
                temp_d = data[i]
            # rows = create_rows(data)
            all_data.append((pull_date.isoformat(), line, data))
            print pull_date.isoformat(), line
        except:
            print "error with url {}".format(url)
    pull_date += timedelta(1)
```

Data Source: CitiBike Trips

```
# Loop over each month
for url in $urls
do
    # Download the zip file
    curl -O $url

    # Define local file names
    file=`basename $url`
    csv=${file//.zip/}".csv"

    # Unzip the downloaded file, remove header lines, include only relevant columns,
    # remove minute and second from timestamp, group/count each unique line, and save as csv
    unzip -p $file | sed 1d | cut -d, -f2,4 | sed 's/:[0-9][0-9]:[0-9][0-9]//g' | sort | uniq -c > $csv


    # Remove the zip file
    rm $file
done
```

Data Source: CitiBike Locations

```
- stationBeanList: [
  - {
    id: 72,
    stationName: "W 52 St & 11 Ave",
    availableDocks: 37,
    totalDocks: 39,
    latitude: 40.76727216,
    longitude: -73.99392888,
    statusValue: "In Service",
    statusKey: 1,
    availableBikes: 2,
    stAddress1: "W 52 St & 11 Ave",
    stAddress2: "",
    city: "",
    postalCode: "",
    location: "",
    altitude: "",
    testStation: false,
    lastCommunicationTime: null,
    landMark: ""
  },
  import urllib
  import json

  url = 'http://www.citibikenyc.com/stations/json/'
  response = urllib.urlopen(url)
  data = json.loads(response.read()).get('stationBeanList')
```


Analyzing the Graph of Citibike Stations



Difference Between Graph Structure Metrics

- Average In Degree, Average Out Degree
- Average Path Length Between Stations
- Eigenvector Centrality (Bonacich, 1972, Journal of Mathematical Sociology)
 - Original basis for Google's PageRank
 - Calculate eigenvector for largest eigenvalue of adjacency matrix
 - Each entry in this vector is a centrality measure for a node
 - Measure is guaranteed unique by Perron-Frobenius theorem
 - All nodes having positive centrality measures requires greatest eigenvalue

$$x_v = \frac{1}{\lambda_{\max}} \sum_{t \in N(v)} x_t = \frac{1}{\lambda_{\max}} \sum_{t \in G} a_{v,t} x_t$$

Simple Metrics - In Degree

Weekday

Ending Station ID	In Degree
473	332
151	331
412	331
497	331
293	330
312	330

Weekend

Ending Station ID	In Degree
151	321
293	320
387	317
312	316
263	315
428	315

Simple Metrics - Out Degree

Weekday

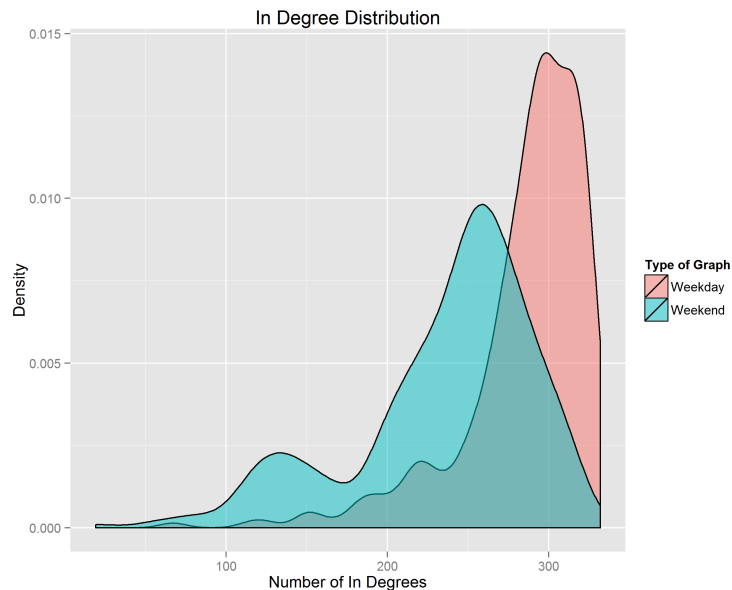
Starting Station ID	Out Degree
151	331
497	331
265	329
387	329
236	328
250	328

Weekend

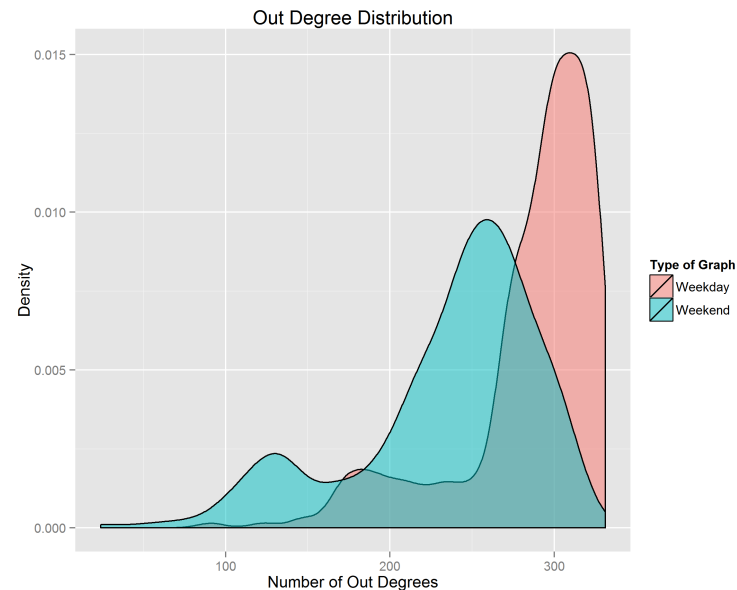
Starting Station ID	Out Degree
312	318
151	316
251	314
237	313
236	312
387	312

Simple Metrics - Degree Distribution


In Degree



Out Degree



Less Simple Metrics - Average Path Length Between Stations



Complicated Metrics - Eigencentrality

Weekday

Station ID	
473	0.0619
151	0.0618
497	0.0618
412	0.0618
293	0.0617
312	0.0617

Weekend

Station ID	Out Degree
151	0.0695
293	0.0695
312	0.0690
250	0.0686
251	0.0686
428	0.0685

Complicated Metrics - Eigencentralities

